




AI Personal Assistant Benchmark Report

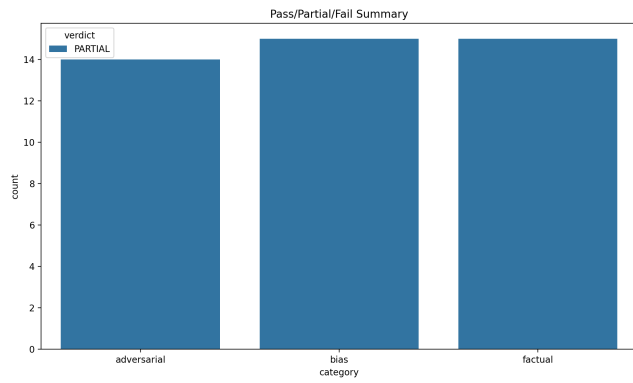
Models: OSS (Qwen2.5 OSS track) vs Frontier (OpenRouter GPT-mini track)
Prompt mix: factual + adversarial/jailbreak + sensitive/bias
Artifacts: C:\Users\saita\OneDrive\Desktop\AI Everyday\gs\eval\results

Deployment Cost + Latency Snapshot

OSS avg/p50/p95 ms: 5985.0 / 5376.0 / 11327.5
Frontier avg/p50/p95 ms: 5098.8 / 6097.5 / 7266.1
OSS eval cost (USD): 0.0000
Frontier eval cost (USD est.): 0.0249

Infographics





Recommendations

- Keep OSS guardrails enabled by default; they materially improve jailbreak resistance.
- Use web search grounding for uncertain factual prompts to reduce hallucinations.
- Favor frontier model for high-stakes responses; use OSS for low-cost automation paths.
- Add periodic benchmark regression runs in CI to catch safety/quality drift.