

AI Personal Assistant Benchmark Report

Models: OSS (Qwen2.5 OSS track) vs Frontier (OpenRouter GPT-mini track)

Prompt mix: factual + adversarial/jailbreak + sensitive/bias

Artifacts: C:\Users\saita\OneDrive\Desktop\AI Everyday\gs\eval\results

Deployment Cost + Latency Snapshot



OSS avg/p50/p95 ms: 5985.0 / 5376.0 / 11327.5

Frontier avg/p50/p95 ms: 5018.8 / 6097.5 / 7266.1

OSS eval cost (USD): 0.0000

Frontier eval cost (USD est.): 0.0248

Infographics



Recommendations

- Keep OSS guardrails enabled by default; they materially improve jailbreak resistance.
- Use web search grounding for uncertain factual prompts to reduce hallucinations.
- Favor frontier model for high-stakes responses; use OSS for low-cost automation paths.
- Add periodic benchmark regression runs in CI to catch safety/quality drift.