




AI Personal Assistant Benchmark Report


Models: OSS (Qwen2.5 OSS track) vs Frontier (OpenAI GPT-4.1 track)
Prompt mix: factual + adversarial/jailbreak + sensitive/bias
Artifacts: C:\Users\saita\OneDrive\Desktop\AI Everyday\gs\eval\results

Deployment Cost + Latency Snapshot

OSS avg/p50/p95 ms 5009.8 / 4800.5 / 6852.3
Frontier avg/p50/p95 ms 3069.4 / 7420.0 / 14516.7
OSS eval cost (USD) 0.0000
Frontier eval cost (USD est.): 0.2416

Infographics





Recommendations

- Keep OSS guardrails enabled by default; they materially improve jailbreak resistance.
- Use web search grounding for uncertain factual prompts to reduce hallucinations.
- Favor frontier model for high-stakes responses; use OSS for low-cost automation paths.
- Add periodic benchmark regression runs in CI to catch safety/quality drift.